

Deep Reinforcement Learning for Electric Vehicle Routing Optimization

Energy-Minimizing Electric Vehicle Routing Problem (EM-EVRP)

Dhia Elhak Ezzeddini Salem Fradi

Higher School of Communication of Tunis
Carthage University


Supervisor: Dr. Ferdaoues Chaabane
Host Company: Sagemcom

Academic Year 2025–2026

Outline

- 1 Introduction
- 2 Problem Statement
- 3 Mathematical Formulation
- 4 Data Generation
- 5 Baseline Optimization Methods
- 6 Deep Reinforcement Learning Approach
- 7 System Architecture
- 8 Results and Evaluation

Challenges in EV Logistics




Electric Vehicle Routing Problem (EVRP)

Classical VRP:

- Serve all customers from depot
- Minimize total travel distance/cost
- Respect vehicle capacity

EVRP Extension:

- **Battery capacity** constraints
- **Charging station** visits allowed
- **Energy consumption** varies with load
- **Recharging time** impacts schedule



Multiple routes with charging stops

Energy-Minimizing EVRP (EM-EVRP)

Problem Definition

Objective: Minimize **total energy consumption** while serving all customers

Subject to Constraints:

① Customer Service

- Each customer visited exactly once
- Full demand satisfaction

② Vehicle Capacity

- Total load $\leq$ max capacity
- Partial deliveries possible

③ Battery SOC

- SOC ≥ 0 at all times
- Recharge at stations

④ Time Windows

- Service within $[a_i, b_i]$
- Route time budget

⑤ Route Feasibility

- Start and end at depot
- Valid path existence

⑥ Charging Logic

- Optional station visits
- Charging time cost

MILP Formulation: Sets and Parameters

Sets

N	All nodes (depot + customers)
C	Customers $C = N \setminus \{0\}$
S	Charging stations
K	Fleet of vehicles
R	Trip indices

Network Parameters

d_i	Demand at customer i
Q_k	Capacity of vehicle k
t_{ij}	Travel time on arc (i, j)
$[a_i, b_i]$	Time window at node i

EV Energy Parameters

m_c	Vehicle curb mass
C_d	Aerodynamic drag coeff.
C_r	Rolling resistance coeff.
ρ	Air density
A	Frontal area
α_{ij}	Slope angle on arc
$\varphi_{d/r}$	Motor efficiency
$\phi_{d/r}$	Battery efficiency

Physics-Based Energy Consumption Model

Total Vehicle Mass on Arc (i, j)

$$m_{ij} = m_c + u_{ij}$$

where u_{ij} is the current cargo load

Mechanical Power Required

$$P_{ij} = \underbrace{\frac{1}{2} C_d \rho A v_{ij}^2}_{\text{Aerodynamic drag}} + \underbrace{m_{ij} g \sin \alpha_{ij}}_{\text{Gravitational}} + \underbrace{m_{ij} g C_r \cos \alpha_{ij}}_{\text{Rolling resistance}}$$

Energy Consumption with Efficiency

$$E_{ij} = \begin{cases} \phi_d \cdot \varphi_d \cdot P_{ij} \cdot t_{ij} & \text{if } P_{ij} \geq 0 \text{ (propulsion)} \\ \phi_r \cdot \varphi_r \cdot P_{ij} \cdot t_{ij} & \text{if } P_{ij} < 0 \text{ (regeneration)} \end{cases}$$

Objective: Minimize Total Energy

$$\min \sum_{k \in K} \sum_{r \in R} \sum_{(i,j) \in A} E_{ij} \cdot x_{ijk}^r$$

Customer Service: $\sum_{k,r} y_{ik}^r = 1$

Capacity: $\sum_i d_i y_{ik}^r \leq Q_k$

Flow: $\sum_j x_{ijk}^r = y_{ik}^r$

MTZ: $w_{ik}^r - w_{jk}^r + |C| x_{ijk}^r \leq |C| - 1$

NP-Hard $\Rightarrow$ Exact methods don't scale


Data Generation Pipeline

Training Data:

- Random node coordinates on 100×100 grid
- Random elevations for slope calculation
- Random customer demands
- Pairwise distances and slopes computed

Dataset Size

Training	1,280,000
Validation	10,000
Batch size	1,024



Classical Optimization Approaches

We implemented **four baseline methods** for comparison:

1. Google OR-Tools

- Industry-standard solver
- Constraint programming
- Arc cost callbacks for energy
- Local search improvements

2. Hill Climbing

- Simple local search
- 2-opt, Or-opt, Exchange
- Fast but local optima


3. Simulated Annealing

- Probabilistic acceptance
- Escapes local optima
- Temperature-based control
- $P(\text{accept}) = e^{-\Delta E/T}$

4. Ant Colony Optimization

- Nature-inspired
- Pheromone trails
- Collective intelligence

Local Search: 2-Opt Operator



2-Opt Improvement

Reverse a segment of the route to eliminate crossing edges and reduce total cost

Simulated Annealing: Acceptance Rule

Probabilistic Acceptance

$$P(\text{accept worse solution}) = \exp\left(-\frac{\Delta E}{T}\right)$$

Temperature Behavior:

- High $T \Rightarrow$ explore (accept worse)
- Low $T \Rightarrow$ exploit (only improvements)


Key Parameters:

- T_{start} : initial temperature
- T_{end} : final temperature
- Cooling: $T_{k+1} = \alpha \cdot T_k$

Challenge

How to choose T_{start} and T_{end} for optimal performance?

SA Parameter Tuning: Initial Exploration



Experiment Setup:

- Monte Carlo sampling
- Vary T_{start} and T_{end}
- Record score distribution

Observation

High variance across parameter space
— need to narrow down!


SA Parameter Tuning: Restricting T_{end}

Restriction Applied:

$$T_{\text{end}} < 0.01$$

Finding

Lower end temperatures lead to better convergence — algorithm should cool down sufficiently




SA Parameter Tuning: Restricting T_{start}

Restrictions Applied:


$$T_{\text{start}} \geq 0.1 \quad \text{and} \quad T_{\text{end}} < 0.01$$

Finding

Starting temperature must be high enough for initial exploration



SA Parameter Tuning: Heatmap Analysis



Methodology:


- 10×10 parameter grid
- 5000 Monte Carlo iterations
- Track best solution frequency

Optimal Region

$$T_{\text{start}} \approx 0.2$$

$$T_{\text{end}} \approx 5 \times 10^{-3}$$

SA Convergence with Optimal Parameters



Configuration

$$T_{\text{start}} = 0.2, \quad T_{\text{end}} = 5 \times 10^{-3}$$

Probabilistic Transition Rule:

$$p_{ij}^a = \frac{[\tau_{ij}]^\alpha \cdot [\eta_{ij}]^\beta}{\sum_{l \in \mathcal{N}_i} [\tau_{il}]^\alpha \cdot [\eta_{il}]^\beta}$$

Pheromone Update:

$$\tau_{ij} \leftarrow (1 - \rho)\tau_{ij} + \sum_a \Delta\tau_{ij}^a$$

Better solutions deposit more pheromone


Parameters (Literature-based)

α	=	3 (pheromone weight)
β	=	1 (heuristic weight)
ρ	=	0.1 (evaporation)
Q	=	5 (deposit constant)
Ants	=	80
Epochs	=	600

Parameter Choice

$\alpha > \beta$: pheromone dominates heuristic
(exploitation over exploration)

ACO: Demo Result



Observation

ACO constructs feasible routes visiting all customers with charging station detours when needed

Why Deep Reinforcement Learning?

Traditional Methods Limitations

- Exact methods: exponential time
- Heuristics: problem-specific design
- No generalization to new instances
- Slow for real-time applications

DRL Advantages

- **Learn** to solve, not program
- **Generalize** to new instances
- **Fast inference** (milliseconds)
- **End-to-end** optimization

Key Insight

Train a neural network to **construct** solutions step-by-step, learning from experience which decisions lead to lower energy consumption

Inspired by: Kool et al., “Attention, Learn to Solve Routing Problems!” (ICLR 2019)

Markov Decision Process Components

State s_t : Current partial solution + vehicle status

- Visited nodes, current position
- Remaining capacity, SOC, time budget

Action a_t : Select next node to visit

- Customer, charging station, or depot

Reward r_t : Negative energy for traversed arc


- Minimize cumulative energy consumption

Policy π_θ : Neural network with parameters θ

- Maps states to action probabilities

$$\textbf{Objective: } \max_{\theta} \mathbb{E}_{\pi_{\theta}} [\sum_t r_t] = \max_{\theta} \mathbb{E}_{\pi_{\theta}} [-\sum_t E_t]$$

Attention-Based Architecture: Overview



Graph Attention Encoder

Multi-Head Self-Attention Layer

$$h_i^{(l)} = \text{BN} \left(h_i^{(l-1)} + \text{MHA} \left(h_i^{(l-1)}, H^{(l-1)} \right) \right)$$

Attention Mechanism:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Graph Embedding:

$$\bar{h} = \frac{1}{n} \sum_{i=1}^n h_i$$

Mean aggregation of all node embeddings

Architecture Details

Embedding dim	128
Attention heads	8
Encoder layers	3
FF hidden	512
Normalization	Batch

Context-Aware Decoder

Context Embedding

$$c_t = W_c [\bar{h}; h_{\pi_{t-1}}; \text{load}_t; \text{SOC}_t; \text{time}_t]$$

Combines: graph embedding + previous node + vehicle state

Attention-Based Selection

$$u_{tj} = \begin{cases} C \cdot \tanh\left(\frac{q_t \cdot k_j}{\sqrt{d_k}}\right) & \text{if } j \in \mathcal{A}_t \\ -\infty & \text{otherwise} \end{cases}$$

$$p_{\theta}(a_t = j | s_t) = \text{softmax}(u_t)_j$$

Feasibility Masking

$\mathcal{A}_t$ = nodes that satisfy: demand $\leq$ capacity AND energy needed $\leq$ SOC

REINFORCE Training Algorithm

Policy Gradient Objective

$$J(\theta) = \mathbb{E}_{\pi_{\theta}} [L(\pi)] \quad \text{where } L(\pi) \text{ is route energy}$$

Gradient with Rollout Baseline

$$\nabla_{\theta} J \approx \frac{1}{B} \sum_{i=1}^B (L(\pi_i) - b(s_i)) \nabla_{\theta} \log p_{\theta}(\pi_i | s_i)$$

$b(s)$ = baseline from greedy policy copy (reduces variance)

Baseline Update

Update when policy shows significant improvement (paired t-test, $p < 0.05$)

State of Charge Tracking

$$\text{SOC}_{t+1} = \text{SOC}_t - \frac{E_{ij}}{E_{\text{battery}}}$$

- Continuous SOC monitoring
- Reset at charging stations
- Negative consumption = regeneration

Load Dynamics

$$\text{load}_{t+1} = \text{load}_t - d_{\text{served}}$$

Cargo decreases after each delivery

Dynamic Masking Rules

- 1 Sufficient SOC for arc?
- 2 Can reach depot from there?
- 3 Enough capacity for demand?
- 4 Time budget remaining?
- 5 Station detour feasible?

Learning Charging Decisions

Model learns **when** to visit charging stations through experience

Training Configuration

Problem Configuration

Customers	10
Charging stations	4
Sequence length	16
Initial SOC	80%
Max load	50 units

Training Parameters

Epochs	20
Batch size	1,024
Training samples	1,280,000
Learning rate	10^{-4}
Baseline	Rollout


Model Architecture

Embedding dim	128
Encoder layers	3
Attention heads	8
FF hidden	512
Tanh clipping	10
Optimizer	Adam

Hardware

GPU	NVIDIA RTX
Training time	~8 hours

Training Results



Observations

Consistent improvement, convergence at epoch 80

Generalization

Validation follows training, no overfitting

Performance Comparison

Method	Energy (kWh)	Inference Time
OR-Tools	11.92	2.1s
DRL	12.18	0.5s
Ant Colony Optimization	12.67	3.2s
Simulated Annealing	12.89	1.5s
Hill Climbing	13.21	0.8s

DRL Performance

- Only **2.2%** gap vs OR-Tools
- **7.8%** better than Hill Climbing
- Near-optimal solution quality

Speed Advantage

- **4x** faster than OR-Tools
- **6x** faster than ACO
- Real-time capable

Open Source Routing Machine (OSRM):

- Production-grade routing engine
- Real street-level distances
- Fast table API for distance matrices
- OpenStreetMap data

Integration Pipeline:

- 1 Upload customer GPS coordinates
- 2 Query OSRM for distance matrix
- 3 Run DRL model inference
- 4 Reconstruct routes on map
- 5 Calculate real energy consumption

Scalability Strategy


K-means Clustering:

- Group customers into batches
- 10 customers per group
- Run DRL on each group
- Combine results

NYC Configuration


50 pre-configured charging stations
across New York City

Interactive Dashboard: Main Interface




Upload customers • Configure EV • Select stations • Real-time optimization

Interactive Dashboard: Route Optimization





Inference Process

- Interactive map with customer locations and charging stations
- Real-time route visualization with street-level paths
- Color-coded routes for multiple vehicles



Optimized Routes

Interactive Dashboard: Analytics



Analytics Features

- Energy consumption breakdown per route segment
- SOC tracking throughout the journey
- Cost analysis and performance metrics
- Export results as JSON for integration

Summary of Contributions

① Comprehensive MILP Formulation

- Physics-based energy consumption model
- Multi-trip capability with time windows
- Regenerative braking consideration

② Baseline Implementation

- OR-Tools, Hill Climbing, SA, ACO
- Standardized benchmarking framework

③ Deep Reinforcement Learning Solution

- Attention-based encoder-decoder
- EV-specific state representation
- Fast and generalizable

④ Production-Ready System

- Interactive Streamlit dashboard
- OSRM integration for real streets
- Export and analytics features

Limitations and Future Work

Current Limitations

- Trained on 10 customers
- Single depot scenario
- Homogeneous fleet
- Static demand
- No traffic dynamics

Future Directions

- Scale to 50+ customers
- Multi-depot extension
- Heterogeneous EV types
- Dynamic reoptimization
- Real-time traffic integration
- Joint routing + charging scheduling

Research Extensions

- Transfer learning for larger instances
- Multi-agent coordination for fleet management
- Integration with smart grid for optimal charging times

Thank You!

Backup: References (1/2)

-  W. Kool, H. van Hoof, M. Welling, *“Attention, Learn to Solve Routing Problems!”*, ICLR, 2019.
-  M. Tang et al., *“Energy-optimal routing for EVs using DRL with transformer,”* Applied Energy, 2023.
-  T. Erdelić, T. Carić, *“A Survey on the EVRP: Variants and Solution Approaches,”* J. Adv. Transportation, 2019.
-  S. Kirkpatrick et al., *“Optimization by Simulated Annealing,”* Science, 1983.
-  M. Dorigo et al., *“The Ant System: Optimization by cooperating agents,”* IEEE Trans. SMC, 1996.
-  G. B. Dantzig, J. H. Ramser, *“The Truck Dispatching Problem,”* Management Science, 1959.

Backup: References (2/2)

-  E. W. Dijkstra, “*A note on two problems in connexion with graphs*,” Numerische Mathematik, 1959.
-  J. MacQueen, “*Some methods for classification and analysis of multivariate observations*,” Berkeley Symp., 1967.
-  A. Vaswani et al., “*Attention Is All You Need*,” NeurIPS, 2017.
-  Google OR-Tools. <https://developers.google.com/optimization>
-  Project OSRM. <http://project-osrm.org/>
-  CVRPLIB. <http://vrp.atd-lab.inf.puc-rio.br/>
-  G. Boeing, “*OSMnx: Street Networks*,” Computers, Environment and Urban Systems, 2017.